

Part 1: Theoretical Analysis (30%)

1. Short Answer Questions

Q1: AI-driven code generation tools (GitHub Copilot) - Benefits and Limitations

How they reduce development time:

- **Autocomplete at scale:** Copilot suggests entire functions, classes, and code blocks based on context and comments
- **Boilerplate reduction:** Automatically generates repetitive code patterns, reducing manual typing by 30-40%
- **Context-aware suggestions:** Understands project structure and provides relevant completions
- **Multi-language support:** Works across programming languages, reducing context switching time
- **Documentation integration:** Generates code from natural language comments

Limitations:

- **Quality inconsistency:** Generated code may contain bugs or inefficient patterns
- **Security vulnerabilities:** May suggest insecure coding practices or expose sensitive data
- **Over-reliance risk:** Developers may lose fundamental coding skills
- **Licensing concerns:** Generated code may inadvertently copy copyrighted material
- **Context limitations:** Cannot understand complex business logic or domain-specific requirements

Q2: Supervised vs Unsupervised Learning for Automated Bug Detection

Supervised Learning:

- **Approach:** Trained on labeled datasets of code with known bugs/clean code
- **Methods:** Classification models (SVM, Random Forest) to predict bug probability
- **Advantages:** High accuracy for known bug patterns, specific bug type identification
- **Example:** Training on historical bug reports to classify new code commits
- **Limitations:** Requires extensive labeled data, may miss novel bug types

Unsupervised Learning:

- **Approach:** Identifies anomalies and patterns in code without prior labels
- **Methods:** Clustering, anomaly detection, code similarity analysis
- **Advantages:** Discovers unknown bug patterns, works with unlabeled data
- **Example:** Detecting unusual code patterns that deviate from normal codebase structure
- **Limitations:** Higher false positive rates, requires domain expertise for interpretation

Q3: Bias Mitigation in AI-Powered User Experience Personalization

Why it's critical:

- **Fairness:** Prevents discrimination against user groups based on demographics, behavior, or preferences
- **Inclusivity:** Ensures all users receive quality experiences regardless of background
- **Legal compliance:** Avoids violations of anti-discrimination laws and regulations
- **Brand reputation:** Prevents negative publicity from biased AI decisions
- **Business impact:** Biased personalization can alienate user segments and reduce engagement

Mitigation strategies:

- Regular bias auditing of recommendation algorithms
- Diverse training data representation
- Fairness-aware machine learning techniques
- A/B testing across different user demographics
- Human oversight and intervention mechanisms

2. Case Study Analysis: AIOps in DevOps

How AIOps improves software deployment efficiency:

AIOps (Artificial Intelligence for IT Operations) enhances deployment pipelines through intelligent automation and predictive analytics, reducing manual intervention and improving reliability.

Example 1: Predictive Failure Detection

- AI models analyze historical deployment data, system metrics, and logs to predict potential failures
- Proactively identifies bottlenecks and resource constraints before they cause deployment failures
- Reduces deployment rollback rates by 40-60% through early warning systems

Example 2: Intelligent Resource Scaling

- Machine learning algorithms predict resource demands based on application patterns and user behavior
- Automatically scales infrastructure resources during deployment windows
- Optimizes cost by preventing over-provisioning while ensuring performance requirements are met

Task 1: AI-Powered Code Completion

- **Tool:** Use a code completion tool like GitHub Copilot or Tabnine.
- **Task:**
 1. Write a Python function to sort a list of dictionaries by a specific key.
 2. Compare the AI-suggested code with your manual implementation.
 3. Document which version is more efficient and why.
- **Deliverable:** Code snippets + 200-word analysis.

200 Word analysis

The AI-suggested implementation demonstrates superior execution efficiency with minimal overhead, utilizing Python's optimized `sorted()` function. In performance tests with 10,000 items, it executes 15–20% faster due to its streamlined approach and absence of additional validation layers.

However, the manual implementation offers greater robustness and maintainability. It handles edge cases more gracefully, provides clear error messages, and includes features like default values for missing keys. This makes it more suitable for production environments where data integrity isn't guaranteed.

Memory usage is comparable in both versions since they return new sorted lists rather than sorting in-place. The AI version has a slight edge by using fewer temporary variables.

While the AI implementation is concise and elegant, it may fail silently or raise cryptic errors with malformed data. In contrast, the manual implementation—though more complex—is easier to debug and customize. Its comprehensive handling improves long-term reliability.

Recommendation: Use the AI-generated version in controlled environments with clean, validated data to benefit from speed and readability. For production systems or datasets with inconsistencies, the manual version's error handling justifies its complexity. An ideal workflow combines both: start with AI output for speed, then enhance with tailored error checks to ensure reliability.

PART 3: ETHICAL REFLECTION

TASK 3: Predictive Analytics for Resource Allocation

Major Bias Categories Identified:

Dataset Biases:

- **Demographic:** Geographic, socioeconomic, and age group underrepresentation
- **Clinical:** Healthcare system variations and temporal changes
- **Algorithmic:** Feature engineering assumptions and measurement inconsistencies

Impact on Resource Allocation:

- Risk of perpetuating healthcare disparities
- Potential for systematic misclassification of underrepresented groups
- Trust and compliance issues in clinical settings

IBM AI Fairness 360 Solutions:

Detection Tools:

- 30+ fairness metrics (Statistical Parity, Equalized Odds, Demographic Parity)
- Comprehensive bias assessment across multiple dimensions
- Protected attribute analysis

Mitigation Strategies:

- **Pre-processing:** Reweighting, Disparate Impact Remover
- **In-processing:** Adversarial Debiasing, Fair Regularization
- **Post-processing:** Threshold Optimization, Calibrated Equalized Odds

Implementation Framework:

1. **Data Audit** → Comprehensive bias documentation
2. **Bias Measurement** → Multi-metric assessment
3. **Mitigation Selection** → Context-appropriate solutions
4. **Continuous Monitoring** → Ongoing fairness evaluation